

HW: X-ray images classification

Before you begin, open Mobaxterm and connect to triton with the user and password you were give with. Activate the environment `2ndPaper` and then type the command `pip install scikit-image`.

In this assignment you will be dealing with classification of 32X32 X-ray images of the chest. The image can be classified into one of four options: lungs (l), clavicles (c), and heart (h) and background (b). Even though those labels are dependent, we will treat this task as multiclass and not as multilabel. The dataset for this assignment is located on a shared folder on triton (`/MLdata/MLcourse/X_ray/`).

```
In [1]: import os
import numpy as np
from tensorflow.keras.layers import Dense, MaxPool2D, Conv2D, Dropout
from tensorflow.keras.layers import Flatten, InputLayer
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import *

from tensorflow.keras.initializers import Constant
from tensorflow.keras.datasets import fashion_mnist
import tensorflow.keras.backend as K
from tensorflow.keras import regularizers
from tensorflow import keras
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import *
from skimage.io import imread

from skimage.transform import rescale, resize, downscale_local_mean
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rc('axes', labelsz=14)
mpl.rc('xtick', labelsz=12)
mpl.rc('ytick', labelsz=12)
os.environ["CUDA_DEVICE_ORDER"]="PCI_BUS_ID"
os.environ["CUDA_VISIBLE_DEVICES"]="2"
```

```
In [2]: import tensorflow as tf
config = tf.compat.v1.ConfigProto(gpu_options =
                                   tf.compat.v1.GPUOptions(per_process_gpu_memory_fraction=0.8)
                                   # device_count = {'GPU': 1}
                                   )
config.gpu_options.allow_growth = True
session = tf.compat.v1.Session(config=config)
tf.compat.v1.keras.backend.set_session(session)
```

```
In [3]: def preprocess(datapath):
        # This part reads the images
        classes = ['b', 'c', 'l', 'h']
        imagelist = [fn for fn in os.listdir(datapath)]
        N = len(imagelist)
        num_classes = len(classes)
        images = np.zeros((N, 32, 32, 1))
        Y = np.zeros((N, num_classes))
        ii=0
        for fn in imagelist:

            src = imread(os.path.join(datapath, fn),1)
            img = resize(src,(32,32),order = 3)

            images[ii,:,:,:] = img
            cc = -1
            for cl in range(len(classes)):
                if fn[-5] == classes[cl]:
                    cc = cl
            Y[ii,cc]=1
            ii += 1

        BaseImages = images
        BaseY = Y
        return BaseImages, BaseY
```

```
In [4]: def preprocess_train_and_val(datapath):
# This part reads the images
classes = ['b','c','l','h']
imagelist = [fn for fn in os.listdir(datapath)]
N = len(imagelist)
num_classes = len(classes)
images = np.zeros((N, 32, 32, 1))
Y = np.zeros((N,num_classes))
ii=0
for fn in imagelist:

    images[ii,:,:,:] = imread(os.path.join(datapath, fn),1)
    cc = -1
    for cl in range(len(classes)):
        if fn[-5] == classes[cl]:
            cc = cl
    Y[ii,cc]=1
    ii += 1

return images, Y
```

```
In [5]: #Loading the data for training and validation:
src_data = '/MLdata/MLcourse/X_ray/'
train_path = src_data + 'train'
val_path = src_data + 'validation'
test_path = src_data + 'test'
BaseX_train , BaseY_train = preprocess_train_and_val(train_path)
BaseX_val , BaseY_val = preprocess_train_and_val(val_path)
X_test, Y_test = preprocess(test_path)
```

```
In [6]: keras.backend.clear_session()
```

PART 1: Fully connected layers

Task 1: *NN with fully connected layers.

Elaborate a NN with 2 hidden fully connected layers with 300, 150 neurons and 4 neurons for classification. Use ReLU activation functions for the hidden layers and He_normal for initialization. Don't forget to flatten your image before feedforward to the first dense layer. Name the model `model_relu` .*

```
In [7]: #-----Impelment your code here:-----
#-----
from tensorflow.keras.initializers import he_normal

model_relu = Sequential(name="model_relu")
model_relu.add(Flatten(input_shape=(32, 32, 1)))
model_relu.add(Dense(300, activation='relu', kernel_initializer=tf.keras.initializers.he_normal()))
model_relu.add(Dense(150))
model_relu.add(Activation('relu'))
model_relu.add(Dense(4))
model_relu.add(Activation('softmax'))
#-----
#-----
```

```
In [8]: model_relu.summary()
```

Model: "model_relu"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 300)	307500
dense_1 (Dense)	(None, 150)	45150
activation (Activation)	(None, 150)	0
dense_2 (Dense)	(None, 4)	604
activation_1 (Activation)	(None, 4)	0
=====		
Total params: 353,254		
Trainable params: 353,254		
Non-trainable params: 0		
=====		

```
In [9]: #Inputs:
input_shape = (32,32,1)
learn_rate = 1e-5
decay = 0
batch_size = 64
epochs = 25

#Define your optimizar parameters:
AdamOpt = Adam(lr=learn_rate,decay=decay)
```

Compile the model with the optimizer above, accuracy metric and adequate loss for multiclass task. Train your model on the training set and evaluate the model on the testing set. Print the accuracy and loss over the testing set.

```
In [10]: #-----Implement your code here:-----  
-----  
model_relu.compile(optimizer=AdamOpt, metrics=['accuracy'], loss='categorical_  
crossentropy')  
history_relu = model_relu.fit(BaseX_train, BaseY_train, batch_size=batch_size,  
epochs=epochs, verbose=2, validation_data=(BaseX_val, BaseY_val))  
results_metrics = model_relu.evaluate(X_test, Y_test, batch_size=batch_size)  
print("Test loss, Test accuracy:", results_metrics)  
#-----  
-----
```

Epoch 1/25
102/102 - 1s - loss: 1.3252 - accuracy: 0.3630 - val_loss: 1.2408 - val_accuracy: 0.5301
Epoch 2/25
102/102 - 0s - loss: 1.1845 - accuracy: 0.5748 - val_loss: 1.1275 - val_accuracy: 0.6065
Epoch 3/25
102/102 - 0s - loss: 1.0932 - accuracy: 0.6345 - val_loss: 1.0524 - val_accuracy: 0.6499
Epoch 4/25
102/102 - 0s - loss: 1.0270 - accuracy: 0.6660 - val_loss: 0.9963 - val_accuracy: 0.6881
Epoch 5/25
102/102 - 0s - loss: 0.9743 - accuracy: 0.6960 - val_loss: 0.9501 - val_accuracy: 0.7037
Epoch 6/25
102/102 - 0s - loss: 0.9317 - accuracy: 0.7102 - val_loss: 0.9134 - val_accuracy: 0.7240
Epoch 7/25
102/102 - 0s - loss: 0.8947 - accuracy: 0.7263 - val_loss: 0.8793 - val_accuracy: 0.7297
Epoch 8/25
102/102 - 0s - loss: 0.8604 - accuracy: 0.7374 - val_loss: 0.8507 - val_accuracy: 0.7373
Epoch 9/25
102/102 - 0s - loss: 0.8306 - accuracy: 0.7481 - val_loss: 0.8225 - val_accuracy: 0.7483
Epoch 10/25
102/102 - 0s - loss: 0.8032 - accuracy: 0.7576 - val_loss: 0.7987 - val_accuracy: 0.7569
Epoch 11/25
102/102 - 0s - loss: 0.7794 - accuracy: 0.7695 - val_loss: 0.7771 - val_accuracy: 0.7616
Epoch 12/25
102/102 - 0s - loss: 0.7567 - accuracy: 0.7723 - val_loss: 0.7570 - val_accuracy: 0.7668
Epoch 13/25
102/102 - 0s - loss: 0.7373 - accuracy: 0.7783 - val_loss: 0.7378 - val_accuracy: 0.7726
Epoch 14/25
102/102 - 0s - loss: 0.7181 - accuracy: 0.7830 - val_loss: 0.7211 - val_accuracy: 0.7824
Epoch 15/25
102/102 - 0s - loss: 0.7001 - accuracy: 0.7892 - val_loss: 0.7050 - val_accuracy: 0.7841
Epoch 16/25
102/102 - 0s - loss: 0.6832 - accuracy: 0.7946 - val_loss: 0.6895 - val_accuracy: 0.7882
Epoch 17/25
102/102 - 0s - loss: 0.6675 - accuracy: 0.7978 - val_loss: 0.6763 - val_accuracy: 0.7911
Epoch 18/25
102/102 - 0s - loss: 0.6522 - accuracy: 0.8032 - val_loss: 0.6652 - val_accuracy: 0.7894
Epoch 19/25
102/102 - 0s - loss: 0.6387 - accuracy: 0.8068 - val_loss: 0.6471 - val_accuracy: 0.7969

```

Epoch 20/25
102/102 - 0s - loss: 0.6244 - accuracy: 0.8099 - val_loss: 0.6362 - val_accu
acy: 0.8009
Epoch 21/25
102/102 - 0s - loss: 0.6123 - accuracy: 0.8125 - val_loss: 0.6238 - val_accu
acy: 0.8032
Epoch 22/25
102/102 - 0s - loss: 0.5999 - accuracy: 0.8177 - val_loss: 0.6131 - val_accu
acy: 0.8084
Epoch 23/25
102/102 - 0s - loss: 0.5886 - accuracy: 0.8193 - val_loss: 0.6059 - val_accu
acy: 0.8061
Epoch 24/25
102/102 - 0s - loss: 0.5773 - accuracy: 0.8253 - val_loss: 0.5909 - val_accu
acy: 0.8113
Epoch 25/25
102/102 - 0s - loss: 0.5660 - accuracy: 0.8272 - val_loss: 0.5827 - val_accu
acy: 0.8160
3/3 [=====] - 0s 4ms/step - loss: 0.7957 - accuracy:
0.6800
Test loss, Test accuracy: [0.7957050204277039, 0.6800000071525574]

```

Task 2: *Activation functions.*

Change the activation functions to LeakyRelu or tanh or sigmoid. Name the new model `new_a_model` . Explain how it can affect the model.*

```

In [11]: #------Impelment your code here:-----
-----
new_a_model = Sequential(name="new_a_model")
new_a_model.add(Flatten(input_shape=(32,32,1)))
new_a_model.add(Dense(300,activation='sigmoid', kernel_initializer=tf.keras.in
itializers.he_normal()))
new_a_model.add(Dense(150))
new_a_model.add(Activation('tanh'))
new_a_model.add(Dense(4))
new_a_model.add(Activation('softmax'))
#-----
-----

```

In [12]: `new_a_model.summary()`

Model: "new_a_model"

Layer (type)	Output Shape	Param #
=====		
flatten_1 (Flatten)	(None, 1024)	0
dense_3 (Dense)	(None, 300)	307500
dense_4 (Dense)	(None, 150)	45150
activation_2 (Activation)	(None, 150)	0
dense_5 (Dense)	(None, 4)	604
activation_3 (Activation)	(None, 4)	0
=====		
Total params: 353,254		
Trainable params: 353,254		
Non-trainable params: 0		

Task 3: *Number of epochs.*

Train the new model using 25 and 40 epochs. What difference does it makes in term of performance? Remember to save the compiled model for having initialized weights for every run as we did in tutorial 12. Evaluate each trained model on the test set*

```
In [13]: #Inputs:
input_shape = (32,32,1)
learn_rate = 1e-5
decay = 0
batch_size = 64
epochs = 25

#Defining the optimizar parameters:
AdamOpt = Adam(lr=learn_rate,decay=decay)
```



```
In [14]: #-----Impelment your code here:-----  
-----  
new_a_model.compile(optimizer=AdamOpt, metrics=['accuracy'], loss='categorical  
_crossentropy')  
  
# saving weights:  
if not("init_weights" in os.listdir()):  
    os.mkdir("init_weights")  
save_dir = "init_weights/"  
model_name = "initial_weights"  
model_path = os.path.join(save_dir, model_name)  
new_a_model.save(model_path)  
  
history_new_a = new_a_model.fit(BaseX_train, BaseY_train, batch_size=batch_size,  
epochs=epochs, verbose=2, validation_data=(BaseX_val, BaseY_val))  
results_metrics = new_a_model.evaluate(X_test, Y_test, batch_size=batch_size)  
print("Test loss, Test accuracy:", results_metrics)  
#-----  
-----
```

INFO:tensorflow:Assets written to: init_weights/initial_weights/assets

Epoch 1/25

102/102 - 1s - loss: 1.3973 - accuracy: 0.3267 - val_loss: 1.3272 - val_accuracy: 0.3918

Epoch 2/25

102/102 - 0s - loss: 1.2963 - accuracy: 0.4282 - val_loss: 1.2640 - val_accuracy: 0.5191

Epoch 3/25

102/102 - 0s - loss: 1.2394 - accuracy: 0.5388 - val_loss: 1.2094 - val_accuracy: 0.5579

Epoch 4/25

102/102 - 0s - loss: 1.1906 - accuracy: 0.5544 - val_loss: 1.1628 - val_accuracy: 0.5660

Epoch 5/25

102/102 - 0s - loss: 1.1500 - accuracy: 0.5635 - val_loss: 1.1249 - val_accuracy: 0.5741

Epoch 6/25

102/102 - 0s - loss: 1.1162 - accuracy: 0.5731 - val_loss: 1.0931 - val_accuracy: 0.5845

Epoch 7/25

102/102 - 0s - loss: 1.0875 - accuracy: 0.5817 - val_loss: 1.0659 - val_accuracy: 0.6030

Epoch 8/25

102/102 - 0s - loss: 1.0633 - accuracy: 0.5965 - val_loss: 1.0440 - val_accuracy: 0.6123

Epoch 9/25

102/102 - 0s - loss: 1.0429 - accuracy: 0.5944 - val_loss: 1.0240 - val_accuracy: 0.6233

Epoch 10/25

102/102 - 0s - loss: 1.0245 - accuracy: 0.6115 - val_loss: 1.0075 - val_accuracy: 0.6186

Epoch 11/25

102/102 - 0s - loss: 1.0081 - accuracy: 0.6155 - val_loss: 0.9936 - val_accuracy: 0.6233

Epoch 12/25

102/102 - 0s - loss: 0.9941 - accuracy: 0.6196 - val_loss: 0.9800 - val_accuracy: 0.6285

Epoch 13/25

102/102 - 0s - loss: 0.9812 - accuracy: 0.6231 - val_loss: 0.9675 - val_accuracy: 0.6429

Epoch 14/25

102/102 - 0s - loss: 0.9698 - accuracy: 0.6299 - val_loss: 0.9571 - val_accuracy: 0.6470

Epoch 15/25

102/102 - 0s - loss: 0.9581 - accuracy: 0.6328 - val_loss: 0.9482 - val_accuracy: 0.6458

Epoch 16/25

102/102 - 0s - loss: 0.9479 - accuracy: 0.6386 - val_loss: 0.9373 - val_accuracy: 0.6557

Epoch 17/25

102/102 - 0s - loss: 0.9375 - accuracy: 0.6433 - val_loss: 0.9283 - val_accuracy: 0.6632

Epoch 18/25

102/102 - 0s - loss: 0.9279 - accuracy: 0.6483 - val_loss: 0.9195 - val_accuracy: 0.6620

Epoch 19/25

102/102 - 0s - loss: 0.9196 - accuracy: 0.6492 - val_loss: 0.9129 - val_accuracy: 0.6620

```

acy: 0.6586
Epoch 20/25
102/102 - 0s - loss: 0.9108 - accuracy: 0.6546 - val_loss: 0.9048 - val_accu
acy: 0.6661
Epoch 21/25
102/102 - 0s - loss: 0.9029 - accuracy: 0.6594 - val_loss: 0.8992 - val_accu
acy: 0.6696
Epoch 22/25
102/102 - 0s - loss: 0.8947 - accuracy: 0.6627 - val_loss: 0.8896 - val_accu
acy: 0.6696
Epoch 23/25
102/102 - 0s - loss: 0.8868 - accuracy: 0.6674 - val_loss: 0.8851 - val_accu
acy: 0.6794
Epoch 24/25
102/102 - 0s - loss: 0.8804 - accuracy: 0.6718 - val_loss: 0.8762 - val_accu
acy: 0.6800
Epoch 25/25
102/102 - 0s - loss: 0.8724 - accuracy: 0.6735 - val_loss: 0.8702 - val_accu
acy: 0.6823
3/3 [=====] - 0s 4ms/step - loss: 0.9879 - accuracy:
0.5714
Test loss, Test accuracy: [0.9878899455070496, 0.5714285969734192]

```

```

In [15]: #Inputs:
input_shape = (32,32,1)
learn_rate = 1e-5
decay = 0
batch_size = 64
epochs = 40

#Defining the optimizar parameters:
AdamOpt = Adam(lr=learn_rate,decay=decay)

```

```
In [16]: #-----Impelment your code here:-----  
-----  
from tensorflow.keras.models import load_model  
  
new_a_model=load_model("init_weights/initial_weights")  
  
history_new_a = new_a_model.fit(BaseX_train, BaseY_train, batch_size=batch_size,  
                                epochs=epochs, verbose=2, validation_data=(BaseX_val, BaseY_val))  
results_metrics = new_a_model.evaluate(X_test, Y_test, batch_size=batch_size)  
print("Test loss, Test accuracy:", results_metrics)  
#-----  
-----
```

Epoch 1/40
102/102 - 1s - loss: 1.3987 - accuracy: 0.3315 - val_loss: 1.3277 - val_accuracy: 0.3918
Epoch 2/40
102/102 - 0s - loss: 1.2967 - accuracy: 0.4416 - val_loss: 1.2642 - val_accuracy: 0.5035
Epoch 3/40
102/102 - 0s - loss: 1.2399 - accuracy: 0.5185 - val_loss: 1.2105 - val_accuracy: 0.5579
Epoch 4/40
102/102 - 0s - loss: 1.1915 - accuracy: 0.5609 - val_loss: 1.1662 - val_accuracy: 0.5677
Epoch 5/40
102/102 - 0s - loss: 1.1517 - accuracy: 0.5615 - val_loss: 1.1269 - val_accuracy: 0.5741
Epoch 6/40
102/102 - 0s - loss: 1.1173 - accuracy: 0.5721 - val_loss: 1.0937 - val_accuracy: 0.5880
Epoch 7/40
102/102 - 0s - loss: 1.0884 - accuracy: 0.5845 - val_loss: 1.0668 - val_accuracy: 0.5903
Epoch 8/40
102/102 - 0s - loss: 1.0635 - accuracy: 0.5891 - val_loss: 1.0446 - val_accuracy: 0.5972
Epoch 9/40
102/102 - 0s - loss: 1.0435 - accuracy: 0.5951 - val_loss: 1.0244 - val_accuracy: 0.6244
Epoch 10/40
102/102 - 0s - loss: 1.0251 - accuracy: 0.6112 - val_loss: 1.0082 - val_accuracy: 0.6221
Epoch 11/40
102/102 - 0s - loss: 1.0087 - accuracy: 0.6137 - val_loss: 0.9931 - val_accuracy: 0.6296
Epoch 12/40
102/102 - 0s - loss: 0.9945 - accuracy: 0.6222 - val_loss: 0.9796 - val_accuracy: 0.6360
Epoch 13/40
102/102 - 0s - loss: 0.9817 - accuracy: 0.6247 - val_loss: 0.9682 - val_accuracy: 0.6337
Epoch 14/40
102/102 - 0s - loss: 0.9691 - accuracy: 0.6282 - val_loss: 0.9571 - val_accuracy: 0.6505
Epoch 15/40
102/102 - 0s - loss: 0.9591 - accuracy: 0.6331 - val_loss: 0.9492 - val_accuracy: 0.6429
Epoch 16/40
102/102 - 0s - loss: 0.9485 - accuracy: 0.6389 - val_loss: 0.9374 - val_accuracy: 0.6562
Epoch 17/40
102/102 - 0s - loss: 0.9381 - accuracy: 0.6424 - val_loss: 0.9287 - val_accuracy: 0.6551
Epoch 18/40
102/102 - 0s - loss: 0.9288 - accuracy: 0.6463 - val_loss: 0.9206 - val_accuracy: 0.6632
Epoch 19/40
102/102 - 0s - loss: 0.9207 - accuracy: 0.6508 - val_loss: 0.9133 - val_accuracy: 0.6528

Epoch 20/40
102/102 - 0s - loss: 0.9120 - accuracy: 0.6557 - val_loss: 0.9048 - val_accuracy: 0.6644
Epoch 21/40
102/102 - 0s - loss: 0.9032 - accuracy: 0.6603 - val_loss: 0.8997 - val_accuracy: 0.6678
Epoch 22/40
102/102 - 0s - loss: 0.8955 - accuracy: 0.6637 - val_loss: 0.8918 - val_accuracy: 0.6713
Epoch 23/40
102/102 - 0s - loss: 0.8888 - accuracy: 0.6671 - val_loss: 0.8839 - val_accuracy: 0.6742
Epoch 24/40
102/102 - 0s - loss: 0.8815 - accuracy: 0.6708 - val_loss: 0.8779 - val_accuracy: 0.6742
Epoch 25/40
102/102 - 0s - loss: 0.8742 - accuracy: 0.6708 - val_loss: 0.8722 - val_accuracy: 0.6817
Epoch 26/40
102/102 - 0s - loss: 0.8678 - accuracy: 0.6783 - val_loss: 0.8658 - val_accuracy: 0.6817
Epoch 27/40
102/102 - 0s - loss: 0.8612 - accuracy: 0.6784 - val_loss: 0.8592 - val_accuracy: 0.6846
Epoch 28/40
102/102 - 0s - loss: 0.8541 - accuracy: 0.6852 - val_loss: 0.8542 - val_accuracy: 0.6840
Epoch 29/40
102/102 - 0s - loss: 0.8479 - accuracy: 0.6864 - val_loss: 0.8478 - val_accuracy: 0.6921
Epoch 30/40
102/102 - 0s - loss: 0.8415 - accuracy: 0.6892 - val_loss: 0.8429 - val_accuracy: 0.6916
Epoch 31/40
102/102 - 0s - loss: 0.8355 - accuracy: 0.6894 - val_loss: 0.8369 - val_accuracy: 0.6956
Epoch 32/40
102/102 - 0s - loss: 0.8304 - accuracy: 0.6922 - val_loss: 0.8319 - val_accuracy: 0.7002
Epoch 33/40
102/102 - 0s - loss: 0.8236 - accuracy: 0.6943 - val_loss: 0.8295 - val_accuracy: 0.7014
Epoch 34/40
102/102 - 0s - loss: 0.8182 - accuracy: 0.6954 - val_loss: 0.8227 - val_accuracy: 0.7025
Epoch 35/40
102/102 - 0s - loss: 0.8130 - accuracy: 0.6969 - val_loss: 0.8182 - val_accuracy: 0.7054
Epoch 36/40
102/102 - 0s - loss: 0.8082 - accuracy: 0.6996 - val_loss: 0.8153 - val_accuracy: 0.7031
Epoch 37/40
102/102 - 0s - loss: 0.8025 - accuracy: 0.6996 - val_loss: 0.8083 - val_accuracy: 0.7049
Epoch 38/40
102/102 - 0s - loss: 0.7981 - accuracy: 0.7030 - val_loss: 0.8051 - val_accuracy: 0.7072

```
Epoch 39/40
102/102 - 0s - loss: 0.7931 - accuracy: 0.7051 - val_loss: 0.8000 - val_accu
acy: 0.7072
Epoch 40/40
102/102 - 0s - loss: 0.7886 - accuracy: 0.7059 - val_loss: 0.7958 - val_accu
acy: 0.7089
3/3 [=====] - 0s 4ms/step - loss: 0.9317 - accuracy:
0.5943
Test loss, Test accuracy: [0.9316574335098267, 0.5942857265472412]
```

Task 4: *Mini-batches.*

Build the `model_relu` again and run it with a batch size of 32 instead of 64. What are the advantages of the mini-batch vs. SGD?*

```
In [17]: keras.backend.clear_session()
```

```
In [18]: #-----Impelment your code here:-----
-----
model_relu = Sequential(name="model_relu")
model_relu.add(Flatten(input_shape=(32, 32, 1)))
model_relu.add(Dense(300,activation='relu', kernel_initializer=tf.keras.initia
lizers.he_normal()))
model_relu.add(Dense(150))
model_relu.add(Activation('relu'))
model_relu.add(Dense(4))
model_relu.add(Activation('softmax'))
#-----
-----
```

```
In [19]: batch_size = 32
epochs = 50

#Define your optimizar parameters:
AdamOpt = Adam(lr=learn_rate,decay=decay)
```

```
In [20]: #-----Implement your code here:-----  
-----  
model_relu.compile(optimizer=AdamOpt,metrics=['accuracy'], loss='categorical_crossentropy')  
history_relu = model_relu.fit(BaseX_train, BaseY_train, batch_size=batch_size,  
epochs=epochs, verbose=2, validation_data=(BaseX_val, BaseY_val))  
y_pred_test = model_relu.predict(X_test)  
results_metrics = model_relu.evaluate(X_test, Y_test, batch_size=batch_size)  
print("Test loss, Test accuracy:", results_metrics)  
#-----  
-----
```


Epoch 1/50
203/203 - 1s - loss: 1.2279 - accuracy: 0.5073 - val_loss: 1.1124 - val_accuracy: 0.6047
Epoch 2/50
203/203 - 1s - loss: 1.0546 - accuracy: 0.6452 - val_loss: 1.0010 - val_accuracy: 0.6898
Epoch 3/50
203/203 - 1s - loss: 0.9674 - accuracy: 0.6954 - val_loss: 0.9309 - val_accuracy: 0.7078
Epoch 4/50
203/203 - 1s - loss: 0.9039 - accuracy: 0.7156 - val_loss: 0.8779 - val_accuracy: 0.7269
Epoch 5/50
203/203 - 1s - loss: 0.8524 - accuracy: 0.7349 - val_loss: 0.8340 - val_accuracy: 0.7407
Epoch 6/50
203/203 - 1s - loss: 0.8093 - accuracy: 0.7504 - val_loss: 0.7973 - val_accuracy: 0.7506
Epoch 7/50
203/203 - 1s - loss: 0.7739 - accuracy: 0.7615 - val_loss: 0.7654 - val_accuracy: 0.7604
Epoch 8/50
203/203 - 1s - loss: 0.7418 - accuracy: 0.7734 - val_loss: 0.7391 - val_accuracy: 0.7691
Epoch 9/50
203/203 - 1s - loss: 0.7126 - accuracy: 0.7804 - val_loss: 0.7100 - val_accuracy: 0.7801
Epoch 10/50
203/203 - 1s - loss: 0.6857 - accuracy: 0.7901 - val_loss: 0.6836 - val_accuracy: 0.7818
Epoch 11/50
203/203 - 1s - loss: 0.6610 - accuracy: 0.7980 - val_loss: 0.6642 - val_accuracy: 0.7928
Epoch 12/50
203/203 - 1s - loss: 0.6382 - accuracy: 0.8061 - val_loss: 0.6460 - val_accuracy: 0.7940
Epoch 13/50
203/203 - 1s - loss: 0.6191 - accuracy: 0.8105 - val_loss: 0.6258 - val_accuracy: 0.8015
Epoch 14/50
203/203 - 1s - loss: 0.6003 - accuracy: 0.8136 - val_loss: 0.6090 - val_accuracy: 0.8073
Epoch 15/50
203/203 - 1s - loss: 0.5827 - accuracy: 0.8182 - val_loss: 0.5926 - val_accuracy: 0.8131
Epoch 16/50
203/203 - 1s - loss: 0.5673 - accuracy: 0.8258 - val_loss: 0.5773 - val_accuracy: 0.8177
Epoch 17/50
203/203 - 1s - loss: 0.5515 - accuracy: 0.8323 - val_loss: 0.5637 - val_accuracy: 0.8235
Epoch 18/50
203/203 - 1s - loss: 0.5380 - accuracy: 0.8341 - val_loss: 0.5523 - val_accuracy: 0.8281
Epoch 19/50
203/203 - 1s - loss: 0.5246 - accuracy: 0.8400 - val_loss: 0.5389 - val_accuracy: 0.8310

Epoch 20/50
203/203 - 1s - loss: 0.5117 - accuracy: 0.8407 - val_loss: 0.5256 - val_accuracy: 0.8368
Epoch 21/50
203/203 - 1s - loss: 0.5001 - accuracy: 0.8434 - val_loss: 0.5174 - val_accuracy: 0.8368
Epoch 22/50
203/203 - 1s - loss: 0.4893 - accuracy: 0.8506 - val_loss: 0.5098 - val_accuracy: 0.8397
Epoch 23/50
203/203 - 1s - loss: 0.4800 - accuracy: 0.8513 - val_loss: 0.4983 - val_accuracy: 0.8438
Epoch 24/50
203/203 - 1s - loss: 0.4699 - accuracy: 0.8534 - val_loss: 0.4873 - val_accuracy: 0.8426
Epoch 25/50
203/203 - 1s - loss: 0.4600 - accuracy: 0.8567 - val_loss: 0.4787 - val_accuracy: 0.8461
Epoch 26/50
203/203 - 1s - loss: 0.4518 - accuracy: 0.8590 - val_loss: 0.4695 - val_accuracy: 0.8490
Epoch 27/50
203/203 - 1s - loss: 0.4433 - accuracy: 0.8607 - val_loss: 0.4648 - val_accuracy: 0.8547
Epoch 28/50
203/203 - 1s - loss: 0.4349 - accuracy: 0.8639 - val_loss: 0.4633 - val_accuracy: 0.8530
Epoch 29/50
203/203 - 1s - loss: 0.4287 - accuracy: 0.8669 - val_loss: 0.4502 - val_accuracy: 0.8565
Epoch 30/50
203/203 - 1s - loss: 0.4222 - accuracy: 0.8653 - val_loss: 0.4450 - val_accuracy: 0.8582
Epoch 31/50
203/203 - 1s - loss: 0.4143 - accuracy: 0.8706 - val_loss: 0.4416 - val_accuracy: 0.8623
Epoch 32/50
203/203 - 1s - loss: 0.4081 - accuracy: 0.8723 - val_loss: 0.4308 - val_accuracy: 0.8611
Epoch 33/50
203/203 - 1s - loss: 0.4013 - accuracy: 0.8715 - val_loss: 0.4268 - val_accuracy: 0.8634
Epoch 34/50
203/203 - 1s - loss: 0.3958 - accuracy: 0.8755 - val_loss: 0.4239 - val_accuracy: 0.8652
Epoch 35/50
203/203 - 1s - loss: 0.3902 - accuracy: 0.8772 - val_loss: 0.4193 - val_accuracy: 0.8709
Epoch 36/50
203/203 - 1s - loss: 0.3843 - accuracy: 0.8786 - val_loss: 0.4121 - val_accuracy: 0.8675
Epoch 37/50
203/203 - 1s - loss: 0.3798 - accuracy: 0.8806 - val_loss: 0.4085 - val_accuracy: 0.8721
Epoch 38/50
203/203 - 1s - loss: 0.3742 - accuracy: 0.8814 - val_loss: 0.4051 - val_accuracy: 0.8704

```

Epoch 39/50
203/203 - 1s - loss: 0.3700 - accuracy: 0.8832 - val_loss: 0.3997 - val_accu
acy: 0.8727
Epoch 40/50
203/203 - 1s - loss: 0.3655 - accuracy: 0.8837 - val_loss: 0.3966 - val_accu
acy: 0.8744
Epoch 41/50
203/203 - 1s - loss: 0.3599 - accuracy: 0.8866 - val_loss: 0.3934 - val_accu
acy: 0.8779
Epoch 42/50
203/203 - 1s - loss: 0.3567 - accuracy: 0.8866 - val_loss: 0.3904 - val_accu
acy: 0.8779
Epoch 43/50
203/203 - 1s - loss: 0.3521 - accuracy: 0.8877 - val_loss: 0.3869 - val_accu
acy: 0.8802
Epoch 44/50
203/203 - 1s - loss: 0.3492 - accuracy: 0.8877 - val_loss: 0.3825 - val_accu
acy: 0.8808
Epoch 45/50
203/203 - 1s - loss: 0.3438 - accuracy: 0.8913 - val_loss: 0.3805 - val_accu
acy: 0.8819
Epoch 46/50
203/203 - 1s - loss: 0.3411 - accuracy: 0.8926 - val_loss: 0.3752 - val_accu
acy: 0.8791
Epoch 47/50
203/203 - 1s - loss: 0.3367 - accuracy: 0.8930 - val_loss: 0.3737 - val_accu
acy: 0.8819
Epoch 48/50
203/203 - 1s - loss: 0.3337 - accuracy: 0.8928 - val_loss: 0.3694 - val_accu
acy: 0.8819
Epoch 49/50
203/203 - 1s - loss: 0.3303 - accuracy: 0.8940 - val_loss: 0.3642 - val_accu
acy: 0.8825
Epoch 50/50
203/203 - 1s - loss: 0.3271 - accuracy: 0.8950 - val_loss: 0.3640 - val_accu
acy: 0.8883
6/6 [=====] - 0s 3ms/step - loss: 0.8569 - accuracy:
0.6629
Test loss, Test accuracy: [0.8569125533103943, 0.6628571152687073]

```

*****Task 4:***** *Batch normalization.*

Build the `new_a_model` again and add batch normalization layers. How does it impact your results?*

In [21]: `keras.backend.clear_session()`

```
In [22]: #-----Impelment your code here:-----  
-----  
new_a_model = Sequential(name="new_a_model")  
new_a_model.add(Flatten(input_shape=(32,32,1)))  
new_a_model.add(Dense(300,activation='sigmoid', kernel_initializer=tf.keras.in  
itializers.he_normal()))  
new_a_model.add(BatchNormalization())  
new_a_model.add(Dense(150))  
new_a_model.add(BatchNormalization())  
new_a_model.add(Activation('tanh'))  
new_a_model.add(Dense(4))  
new_a_model.add(BatchNormalization())  
new_a_model.add(Activation('softmax'))  
#-----  
-----
```

```
In [23]: batch_size = 32  
epochs = 50  
  
#Define your optimizar parameters:  
AdamOpt = Adam(lr=learn_rate,decay=decay)  
#Compile the network:
```

```
In [24]: #Performing the training by using fit
#-----Impelment your code here:-----
-----
new_a_model.compile(optimizer=AdamOpt,metrics=['accuracy'], loss='categorical_
crossentropy')
history_new_a = new_a_model.fit(BaseX_train, BaseY_train, batch_size=batch_size, epochs=epochs, verbose=2, validation_data=(BaseX_val, BaseY_val))
y_pred_test = new_a_model.predict(X_test)
results_metrics = new_a_model.evaluate(X_test, Y_test, batch_size=batch_size)
print("Test loss, Test accuracy:", results_metrics)
#-----
-----
```

Epoch 1/50
203/203 - 2s - loss: 1.1146 - accuracy: 0.5431 - val_loss: 1.2678 - val_accuracy: 0.3171
Epoch 2/50
203/203 - 1s - loss: 0.8384 - accuracy: 0.7047 - val_loss: 0.9787 - val_accuracy: 0.6713
Epoch 3/50
203/203 - 1s - loss: 0.7532 - accuracy: 0.7569 - val_loss: 0.7656 - val_accuracy: 0.7691
Epoch 4/50
203/203 - 1s - loss: 0.7073 - accuracy: 0.7830 - val_loss: 0.6745 - val_accuracy: 0.8194
Epoch 5/50
203/203 - 1s - loss: 0.6782 - accuracy: 0.8001 - val_loss: 0.6482 - val_accuracy: 0.8171
Epoch 6/50
203/203 - 1s - loss: 0.6440 - accuracy: 0.8202 - val_loss: 0.6368 - val_accuracy: 0.8322
Epoch 7/50
203/203 - 1s - loss: 0.6292 - accuracy: 0.8261 - val_loss: 0.6358 - val_accuracy: 0.8385
Epoch 8/50
203/203 - 1s - loss: 0.6110 - accuracy: 0.8411 - val_loss: 0.5953 - val_accuracy: 0.8547
Epoch 9/50
203/203 - 1s - loss: 0.5981 - accuracy: 0.8471 - val_loss: 0.5850 - val_accuracy: 0.8588
Epoch 10/50
203/203 - 1s - loss: 0.5865 - accuracy: 0.8545 - val_loss: 0.5821 - val_accuracy: 0.8594
Epoch 11/50
203/203 - 1s - loss: 0.5697 - accuracy: 0.8647 - val_loss: 0.5663 - val_accuracy: 0.8657
Epoch 12/50
203/203 - 1s - loss: 0.5665 - accuracy: 0.8692 - val_loss: 0.5548 - val_accuracy: 0.8704
Epoch 13/50
203/203 - 1s - loss: 0.5567 - accuracy: 0.8712 - val_loss: 0.5525 - val_accuracy: 0.8727
Epoch 14/50
203/203 - 1s - loss: 0.5491 - accuracy: 0.8781 - val_loss: 0.5445 - val_accuracy: 0.8733
Epoch 15/50
203/203 - 1s - loss: 0.5364 - accuracy: 0.8798 - val_loss: 0.5501 - val_accuracy: 0.8744
Epoch 16/50
203/203 - 1s - loss: 0.5338 - accuracy: 0.8826 - val_loss: 0.5416 - val_accuracy: 0.8802
Epoch 17/50
203/203 - 1s - loss: 0.5303 - accuracy: 0.8846 - val_loss: 0.5433 - val_accuracy: 0.8773
Epoch 18/50
203/203 - 1s - loss: 0.5212 - accuracy: 0.8926 - val_loss: 0.5322 - val_accuracy: 0.8779
Epoch 19/50
203/203 - 1s - loss: 0.5166 - accuracy: 0.8892 - val_loss: 0.5128 - val_accuracy: 0.8900

Epoch 20/50
203/203 - 1s - loss: 0.5153 - accuracy: 0.8917 - val_loss: 0.5332 - val_accuracy: 0.8744
Epoch 21/50
203/203 - 1s - loss: 0.5083 - accuracy: 0.8984 - val_loss: 0.5122 - val_accuracy: 0.8918
Epoch 22/50
203/203 - 1s - loss: 0.5017 - accuracy: 0.9042 - val_loss: 0.5040 - val_accuracy: 0.8924
Epoch 23/50
203/203 - 1s - loss: 0.4935 - accuracy: 0.9024 - val_loss: 0.5075 - val_accuracy: 0.8866
Epoch 24/50
203/203 - 1s - loss: 0.4928 - accuracy: 0.9033 - val_loss: 0.5027 - val_accuracy: 0.8958
Epoch 25/50
203/203 - 1s - loss: 0.4904 - accuracy: 0.9033 - val_loss: 0.5013 - val_accuracy: 0.8941
Epoch 26/50
203/203 - 1s - loss: 0.4814 - accuracy: 0.9061 - val_loss: 0.5011 - val_accuracy: 0.8964
Epoch 27/50
203/203 - 1s - loss: 0.4804 - accuracy: 0.9096 - val_loss: 0.4856 - val_accuracy: 0.8993
Epoch 28/50
203/203 - 1s - loss: 0.4734 - accuracy: 0.9101 - val_loss: 0.4880 - val_accuracy: 0.8987
Epoch 29/50
203/203 - 1s - loss: 0.4812 - accuracy: 0.9055 - val_loss: 0.4870 - val_accuracy: 0.8999
Epoch 30/50
203/203 - 1s - loss: 0.4706 - accuracy: 0.9132 - val_loss: 0.4903 - val_accuracy: 0.9005
Epoch 31/50
203/203 - 1s - loss: 0.4687 - accuracy: 0.9135 - val_loss: 0.4772 - val_accuracy: 0.9010
Epoch 32/50
203/203 - 1s - loss: 0.4693 - accuracy: 0.9124 - val_loss: 0.4935 - val_accuracy: 0.8941
Epoch 33/50
203/203 - 1s - loss: 0.4603 - accuracy: 0.9171 - val_loss: 0.4789 - val_accuracy: 0.9051
Epoch 34/50
203/203 - 1s - loss: 0.4586 - accuracy: 0.9116 - val_loss: 0.4789 - val_accuracy: 0.8981
Epoch 35/50
203/203 - 1s - loss: 0.4540 - accuracy: 0.9189 - val_loss: 0.4678 - val_accuracy: 0.9051
Epoch 36/50
203/203 - 1s - loss: 0.4614 - accuracy: 0.9124 - val_loss: 0.4707 - val_accuracy: 0.9045
Epoch 37/50
203/203 - 1s - loss: 0.4486 - accuracy: 0.9240 - val_loss: 0.4662 - val_accuracy: 0.9010
Epoch 38/50
203/203 - 1s - loss: 0.4464 - accuracy: 0.9266 - val_loss: 0.4666 - val_accuracy: 0.9045

```
Epoch 39/50
203/203 - 1s - loss: 0.4440 - accuracy: 0.9209 - val_loss: 0.4593 - val_accu
acy: 0.9045
Epoch 40/50
203/203 - 1s - loss: 0.4383 - accuracy: 0.9300 - val_loss: 0.4553 - val_accu
acy: 0.9068
Epoch 41/50
203/203 - 1s - loss: 0.4389 - accuracy: 0.9291 - val_loss: 0.4627 - val_accu
acy: 0.9045
Epoch 42/50
203/203 - 1s - loss: 0.4359 - accuracy: 0.9262 - val_loss: 0.4761 - val_accu
acy: 0.8918
Epoch 43/50
203/203 - 1s - loss: 0.4420 - accuracy: 0.9240 - val_loss: 0.4512 - val_accu
acy: 0.9086
Epoch 44/50
203/203 - 1s - loss: 0.4348 - accuracy: 0.9243 - val_loss: 0.4501 - val_accu
acy: 0.9080
Epoch 45/50
203/203 - 1s - loss: 0.4301 - accuracy: 0.9282 - val_loss: 0.4569 - val_accu
acy: 0.9109
Epoch 46/50
203/203 - 1s - loss: 0.4313 - accuracy: 0.9257 - val_loss: 0.4471 - val_accu
acy: 0.9126
Epoch 47/50
203/203 - 1s - loss: 0.4256 - accuracy: 0.9311 - val_loss: 0.4520 - val_accu
acy: 0.9080
Epoch 48/50
203/203 - 1s - loss: 0.4260 - accuracy: 0.9297 - val_loss: 0.4494 - val_accu
acy: 0.9115
Epoch 49/50
203/203 - 1s - loss: 0.4230 - accuracy: 0.9320 - val_loss: 0.4610 - val_accu
acy: 0.9062
Epoch 50/50
203/203 - 1s - loss: 0.4182 - accuracy: 0.9320 - val_loss: 0.4477 - val_accu
acy: 0.9120
6/6 [=====] - 0s 3ms/step - loss: 0.9994 - accuracy:
0.5943
Test loss, Test accuracy: [0.9993738532066345, 0.5942857265472412]
```

PART 2: Convolutional Neural Network (CNN)

Task 1: 2D CNN.

Have a look at the model below and answer the following:

- How many layers does it have?
 - How many filter in each layer?
 - Would the number of parmaters be similar to a fully connected NN?
 - Is this specific NN performing regularization?
-

```

In [25]: def get_net(input_shape, drop, dropRate, reg):
    #Defining the network architecture:
    model = Sequential()
    model.add(Permute((1,2,3), input_shape = input_shape))
    model.add(Conv2D(filters=64, kernel_size=(3,3), padding='same', activation
    = 'relu', name='Conv2D_1', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
        model.add(BatchNormalization(axis=1))
        model.add(MaxPooling2D(pool_size=(2, 2)))
        model.add(Conv2D(filters=128, kernel_size=(3,3), padding='same', activation
    = 'relu', name='Conv2D_2', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
        model.add(BatchNormalization(axis=1))
        model.add(Conv2D(filters=128, kernel_size=(3,3), padding='same', activation
    = 'relu', name='Conv2D_3', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
        model.add(BatchNormalization(axis=1))
        model.add(MaxPooling2D(pool_size=(2, 2)))
        model.add(Conv2D(filters=256, kernel_size=(3,3), padding='same', activation
    = 'relu', name='Conv2D_4', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
        model.add(BatchNormalization(axis=1))
        model.add(Conv2D(filters=256, kernel_size=(3,3), padding='same', activation
    = 'relu', name='Conv2D_5', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
        model.add(BatchNormalization(axis=1))
        model.add(MaxPooling2D(pool_size=(2, 2)))
        model.add(Flatten())
    #Fully connected network tail:
    model.add(Dense(512, activation='elu', name='FCN_1'))
    if drop:
        model.add(Dropout(rate=dropRate))
        model.add(Dense(128, activation='elu', name='FCN_2'))
        model.add(Dense(4, activation='softmax', name='FCN_3'))
    model.summary()
    return model

```

```
In [26]: input_shape = (32,32,1)
learn_rate = 1e-5
decay = 1e-03
batch_size = 64
epochs = 25
drop = True
dropRate = 0.3
reg = 1e-2
NNet = get_net(input_shape,drop,dropRate,reg)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
permute (Permute)	(None, 32, 32, 1)	0
Conv2D_1 (Conv2D)	(None, 32, 32, 64)	640
dropout (Dropout)	(None, 32, 32, 64)	0
batch_normalization_3 (Batch Normalization)	(None, 32, 32, 64)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
Conv2D_2 (Conv2D)	(None, 16, 16, 128)	73856
dropout_1 (Dropout)	(None, 16, 16, 128)	0
batch_normalization_4 (Batch Normalization)	(None, 16, 16, 128)	64
Conv2D_3 (Conv2D)	(None, 16, 16, 128)	147584
dropout_2 (Dropout)	(None, 16, 16, 128)	0
batch_normalization_5 (Batch Normalization)	(None, 16, 16, 128)	64
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 128)	0
Conv2D_4 (Conv2D)	(None, 8, 8, 256)	295168
dropout_3 (Dropout)	(None, 8, 8, 256)	0
batch_normalization_6 (Batch Normalization)	(None, 8, 8, 256)	32
Conv2D_5 (Conv2D)	(None, 8, 8, 256)	590080
dropout_4 (Dropout)	(None, 8, 8, 256)	0
batch_normalization_7 (Batch Normalization)	(None, 8, 8, 256)	32
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 256)	0
flatten_1 (Flatten)	(None, 4096)	0
FCN_1 (Dense)	(None, 512)	2097664
dropout_5 (Dropout)	(None, 512)	0
FCN_2 (Dense)	(None, 128)	65664
FCN_3 (Dense)	(None, 4)	516
Total params: 3,271,492		
Trainable params: 3,271,332		
Non-trainable params: 160		

In [27]: `NNet=get_net(input_shape,drop,dropRate,reg)`

Model: "sequential_1"

Layer (type)	Output Shape	Param #
permute_1 (Permute)	(None, 32, 32, 1)	0
Conv2D_1 (Conv2D)	(None, 32, 32, 64)	640
dropout_6 (Dropout)	(None, 32, 32, 64)	0
batch_normalization_8 (Batch Normalization)	(None, 32, 32, 64)	128
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 64)	0
Conv2D_2 (Conv2D)	(None, 16, 16, 128)	73856
dropout_7 (Dropout)	(None, 16, 16, 128)	0
batch_normalization_9 (Batch Normalization)	(None, 16, 16, 128)	64
Conv2D_3 (Conv2D)	(None, 16, 16, 128)	147584
dropout_8 (Dropout)	(None, 16, 16, 128)	0
batch_normalization_10 (Batch Normalization)	(None, 16, 16, 128)	64
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 128)	0
Conv2D_4 (Conv2D)	(None, 8, 8, 256)	295168
dropout_9 (Dropout)	(None, 8, 8, 256)	0
batch_normalization_11 (Batch Normalization)	(None, 8, 8, 256)	32
Conv2D_5 (Conv2D)	(None, 8, 8, 256)	590080
dropout_10 (Dropout)	(None, 8, 8, 256)	0
batch_normalization_12 (Batch Normalization)	(None, 8, 8, 256)	32
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 256)	0
flatten_2 (Flatten)	(None, 4096)	0
FCN_1 (Dense)	(None, 512)	2097664
dropout_11 (Dropout)	(None, 512)	0
FCN_2 (Dense)	(None, 128)	65664
FCN_3 (Dense)	(None, 4)	516
Total params: 3,271,492		
Trainable params: 3,271,332		
Non-trainable params: 160		

```
In [28]: from tensorflow.keras.optimizers import *
import os
from tensorflow.keras.callbacks import *

#Defining the optimizar parameters:
AdamOpt = Adam(lr=learn_rate,decay=decay)

#Compile the network:
NNet.compile(optimizer=AdamOpt, metrics=['acc'], loss='categorical_crossentropy')

#Saving checkpoints during training:
Checkpath = os.getcwd()
Checkp = ModelCheckpoint(Checkpath, monitor='val_acc', verbose=1, save_best_only=True, save_weights_only=True, save_freq=1)
```

```
In [29]: #Performing the training by using fit  
# IMPORTANT NOTE: This will take a few minutes!  
h = NNet.fit(x=BaseX_train, y=BaseY_train, batch_size=batch_size, epochs=epochs,  
            verbose=1, validation_split=0, validation_data = (BaseX_val, BaseY_val), shuffle=True)  
#NNet.save(model_fn)
```


Epoch 1/25
102/102 [=====] - 18s 162ms/step - loss: 8.5007 - acc: 0.3504 - val_loss: 7.9215 - val_acc: 0.2500

Epoch 2/25
102/102 [=====] - 17s 163ms/step - loss: 7.7052 - acc: 0.5118 - val_loss: 8.0498 - val_acc: 0.2512

Epoch 3/25
102/102 [=====] - 17s 165ms/step - loss: 7.5296 - acc: 0.5614 - val_loss: 8.0904 - val_acc: 0.2569

Epoch 4/25
102/102 [=====] - 17s 164ms/step - loss: 7.4107 - acc: 0.5819 - val_loss: 8.1286 - val_acc: 0.2714

Epoch 5/25
102/102 [=====] - 16s 159ms/step - loss: 7.2934 - acc: 0.6304 - val_loss: 8.0654 - val_acc: 0.3056

Epoch 6/25
102/102 [=====] - 16s 159ms/step - loss: 7.1918 - acc: 0.6543 - val_loss: 7.8972 - val_acc: 0.3345

Epoch 7/25
102/102 [=====] - 16s 157ms/step - loss: 7.1311 - acc: 0.6744 - val_loss: 7.8012 - val_acc: 0.3640

Epoch 8/25
102/102 [=====] - 16s 159ms/step - loss: 7.0946 - acc: 0.6798 - val_loss: 7.7292 - val_acc: 0.4005

Epoch 9/25
102/102 [=====] - 16s 157ms/step - loss: 7.0324 - acc: 0.7029 - val_loss: 7.6802 - val_acc: 0.4062

Epoch 10/25
102/102 [=====] - 16s 158ms/step - loss: 6.9823 - acc: 0.7114 - val_loss: 7.7079 - val_acc: 0.4051

Epoch 11/25
102/102 [=====] - 16s 158ms/step - loss: 6.9345 - acc: 0.7223 - val_loss: 7.6792 - val_acc: 0.4051

Epoch 12/25
102/102 [=====] - 16s 156ms/step - loss: 6.8930 - acc: 0.7309 - val_loss: 7.6453 - val_acc: 0.4120

Epoch 13/25
102/102 [=====] - 16s 157ms/step - loss: 6.8857 - acc: 0.7260 - val_loss: 7.6553 - val_acc: 0.3999

Epoch 14/25
102/102 [=====] - 16s 156ms/step - loss: 6.8302 - acc: 0.7549 - val_loss: 7.6467 - val_acc: 0.3987

Epoch 15/25
102/102 [=====] - 16s 157ms/step - loss: 6.7913 - acc: 0.7571 - val_loss: 7.6402 - val_acc: 0.3762

Epoch 16/25
102/102 [=====] - 16s 159ms/step - loss: 6.7759 - acc: 0.7583 - val_loss: 7.6077 - val_acc: 0.3889

Epoch 17/25
102/102 [=====] - 16s 157ms/step - loss: 6.7342 - acc: 0.7663 - val_loss: 7.6194 - val_acc: 0.3808

Epoch 18/25
102/102 [=====] - 16s 157ms/step - loss: 6.6998 - acc: 0.7711 - val_loss: 7.5811 - val_acc: 0.3860

Epoch 19/25
102/102 [=====] - 16s 157ms/step - loss: 6.6817 - acc: 0.7682 - val_loss: 7.5979 - val_acc: 0.3773

```

Epoch 20/25
102/102 [=====] - 16s 156ms/step - loss: 6.6465 - ac
c: 0.7861 - val_loss: 7.6038 - val_acc: 0.3791
Epoch 21/25
102/102 [=====] - 16s 156ms/step - loss: 6.6231 - ac
c: 0.7832 - val_loss: 7.5938 - val_acc: 0.3785
Epoch 22/25
102/102 [=====] - 16s 159ms/step - loss: 6.5786 - ac
c: 0.7937 - val_loss: 7.6143 - val_acc: 0.3727
Epoch 23/25
102/102 [=====] - 16s 157ms/step - loss: 6.5530 - ac
c: 0.8055 - val_loss: 7.5845 - val_acc: 0.3762
Epoch 24/25
102/102 [=====] - 16s 157ms/step - loss: 6.5499 - ac
c: 0.8004 - val_loss: 7.5682 - val_acc: 0.3773
Epoch 25/25
102/102 [=====] - 16s 159ms/step - loss: 6.4970 - ac
c: 0.8132 - val_loss: 7.5504 - val_acc: 0.3808

```

```
In [30]: # NNet.load_weights('Weights_1.h5')
```

```
In [31]: results = NNet.evaluate(X_test,Y_test)
print('test loss, test acc:', results)
```

```

6/6 [=====] - 0s 18ms/step - loss: 7.9150 - acc: 0.3
371
test loss, test acc: [7.9150285720825195, 0.33714285492897034]

```

Task 2: *Number of filters*

Rebuild the function `get_net` to have as an input argument a list of number of filters in each layers, i.e. for the CNN defined above the input should have been `[64, 128, 128, 256, 256]`. Now train the model with the number of filters reduced by half. What were the results.

```

In [32]: #-----Impelment your code here:-----
-----
def get_net(input_shape, drop, dropRate, reg, filters):
    #Defining the network architecture:
    model = Sequential()
    model.add(Permute((1, 2, 3), input_shape=input_shape))
    model.add(Conv2D(filters=filters[0], kernel_size=(3, 3), padding='same', a
ctivation='relu', name='Conv2D_1', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
    model.add(BatchNormalization(axis=1))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Conv2D(filters=filters[1], kernel_size=(3, 3), padding='same', a
ctivation='relu', name='Conv2D_2', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
    model.add(BatchNormalization(axis=1))
    model.add(Conv2D(filters=filters[2], kernel_size=(3, 3), padding='same', a
ctivation='relu', name='Conv2D_3', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
    model.add(BatchNormalization(axis=1))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Conv2D(filters=filters[3], kernel_size=(3, 3), padding='same', a
ctivation='relu', name='Conv2D_4', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
    model.add(BatchNormalization(axis=1))
    model.add(Conv2D(filters=filters[4], kernel_size=(3, 3), padding='same', a
ctivation='relu', name='Conv2D_5', kernel_regularizer=regularizers.l2(reg)))
    if drop:
        model.add(Dropout(rate=dropRate))
    model.add(BatchNormalization(axis=1))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Flatten())
    #Fully connected network tail:
    model.add(Dense(512, activation='elu', name='FCN_1'))
    if drop:
        model.add(Dropout(rate=dropRate))
    model.add(Dense(128, activation='elu', name='FCN_2'))
    model.add(Dense(4, activation='softmax', name='FCN_3'))
    model.summary()
    return model

filters = [32, 64, 64, 128, 128]
NNet_new=get_net(input_shape,drop,dropRate,reg,filters)
NNet_new.compile(optimizer=AdamOpt, metrics=['acc'], loss='categorical_crossen
tropy')
h_new = NNet_new.fit(x=BaseX_train, y=BaseY_train, batch_size=batch_size, epoc
hs=epochs, verbose=1, validation_split=0, validation_data=(BaseX_val, BaseY_va
l), shuffle=True)
results_metrics = NNet_new.evaluate(X_test,Y_test)
print('Test loss, Test acc:', results_metrics)
#-----
-----

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
permute_2 (Permute)	(None, 32, 32, 1)	0
Conv2D_1 (Conv2D)	(None, 32, 32, 32)	320
dropout_12 (Dropout)	(None, 32, 32, 32)	0
batch_normalization_13 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d_6 (MaxPooling2D)	(None, 16, 16, 32)	0
Conv2D_2 (Conv2D)	(None, 16, 16, 64)	18496
dropout_13 (Dropout)	(None, 16, 16, 64)	0
batch_normalization_14 (Batch Normalization)	(None, 16, 16, 64)	64
Conv2D_3 (Conv2D)	(None, 16, 16, 64)	36928
dropout_14 (Dropout)	(None, 16, 16, 64)	0
batch_normalization_15 (Batch Normalization)	(None, 16, 16, 64)	64
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 64)	0
Conv2D_4 (Conv2D)	(None, 8, 8, 128)	73856
dropout_15 (Dropout)	(None, 8, 8, 128)	0
batch_normalization_16 (Batch Normalization)	(None, 8, 8, 128)	32
Conv2D_5 (Conv2D)	(None, 8, 8, 128)	147584
dropout_16 (Dropout)	(None, 8, 8, 128)	0
batch_normalization_17 (Batch Normalization)	(None, 8, 8, 128)	32
max_pooling2d_8 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_3 (Flatten)	(None, 2048)	0
FCN_1 (Dense)	(None, 512)	1049088
dropout_17 (Dropout)	(None, 512)	0
FCN_2 (Dense)	(None, 128)	65664
FCN_3 (Dense)	(None, 4)	516
Total params: 1,392,772		
Trainable params: 1,392,612		
Non-trainable params: 160		

Epoch 1/25

```
102/102 [=====] - 9s 78ms/step - loss: 4.9187 - acc:
0.3881 - val_loss: 4.6354 - val_acc: 0.2500
Epoch 2/25
102/102 [=====] - 8s 75ms/step - loss: 4.4907 - acc:
0.5206 - val_loss: 4.6999 - val_acc: 0.2564
Epoch 3/25
102/102 [=====] - 8s 76ms/step - loss: 4.3440 - acc:
0.5581 - val_loss: 4.7265 - val_acc: 0.2512
Epoch 4/25
102/102 [=====] - 8s 76ms/step - loss: 4.3172 - acc:
0.5594 - val_loss: 4.7363 - val_acc: 0.2564
Epoch 5/25
102/102 [=====] - 8s 76ms/step - loss: 4.2809 - acc:
0.5804 - val_loss: 4.6924 - val_acc: 0.2506
Epoch 6/25
102/102 [=====] - 8s 76ms/step - loss: 4.2296 - acc:
0.5960 - val_loss: 4.6012 - val_acc: 0.3050
Epoch 7/25
102/102 [=====] - 8s 76ms/step - loss: 4.1942 - acc:
0.6107 - val_loss: 4.5395 - val_acc: 0.3461
Epoch 8/25
102/102 [=====] - 8s 77ms/step - loss: 4.1713 - acc:
0.6168 - val_loss: 4.4959 - val_acc: 0.3605
Epoch 9/25
102/102 [=====] - 8s 76ms/step - loss: 4.1553 - acc:
0.6185 - val_loss: 4.4879 - val_acc: 0.3634
Epoch 10/25
102/102 [=====] - 8s 76ms/step - loss: 4.1295 - acc:
0.6235 - val_loss: 4.4830 - val_acc: 0.3628
Epoch 11/25
102/102 [=====] - 8s 76ms/step - loss: 4.0983 - acc:
0.6329 - val_loss: 4.4814 - val_acc: 0.3605
Epoch 12/25
102/102 [=====] - 8s 76ms/step - loss: 4.0771 - acc:
0.6359 - val_loss: 4.4815 - val_acc: 0.3623
Epoch 13/25
102/102 [=====] - 8s 78ms/step - loss: 4.0548 - acc:
0.6397 - val_loss: 4.4844 - val_acc: 0.3617
Epoch 14/25
102/102 [=====] - 8s 76ms/step - loss: 4.0624 - acc:
0.6422 - val_loss: 4.4786 - val_acc: 0.3634
Epoch 15/25
102/102 [=====] - 8s 77ms/step - loss: 4.0375 - acc:
0.6495 - val_loss: 4.4947 - val_acc: 0.3588
Epoch 16/25
102/102 [=====] - 8s 77ms/step - loss: 4.0205 - acc:
0.6460 - val_loss: 4.4989 - val_acc: 0.3524
Epoch 17/25
102/102 [=====] - 8s 77ms/step - loss: 4.0045 - acc:
0.6662 - val_loss: 4.5160 - val_acc: 0.3466
Epoch 18/25
102/102 [=====] - 8s 77ms/step - loss: 3.9997 - acc:
0.6587 - val_loss: 4.5171 - val_acc: 0.3461
Epoch 19/25
102/102 [=====] - 8s 78ms/step - loss: 3.9774 - acc:
0.6745 - val_loss: 4.5293 - val_acc: 0.3414
Epoch 20/25
```

```
102/102 [=====] - 8s 77ms/step - loss: 3.9584 - acc:
0.6731 - val_loss: 4.5325 - val_acc: 0.3426
Epoch 21/25
102/102 [=====] - 8s 77ms/step - loss: 3.9694 - acc:
0.6664 - val_loss: 4.5348 - val_acc: 0.3420
Epoch 22/25
102/102 [=====] - 8s 77ms/step - loss: 3.9771 - acc:
0.6802 - val_loss: 4.5379 - val_acc: 0.3426
Epoch 23/25
102/102 [=====] - 8s 77ms/step - loss: 3.9568 - acc:
0.6741 - val_loss: 4.5434 - val_acc: 0.3409
Epoch 24/25
102/102 [=====] - 8s 78ms/step - loss: 3.9172 - acc:
0.6923 - val_loss: 4.5483 - val_acc: 0.3391
Epoch 25/25
102/102 [=====] - 8s 77ms/step - loss: 3.9014 - acc:
0.6944 - val_loss: 4.5590 - val_acc: 0.3374
6/6 [=====] - 0s 11ms/step - loss: 4.9387 - acc: 0.2
343
Test loss, Test acc: [4.9386749267578125, 0.23428571224212646]
```

That's all folks! See you :)